Stamp-it: A more Thread-efficient, Concurrent Memory Reclamation Scheme in the C++ Memory Model

Manuel Pöter
TU Wien, Faculty of Informatics
Vienna, Austria
manuel@manuel-poeter.at

Jesper Larsson Träff
TU Wien, Faculty of Informatics
Vienna, Austria
traff@par.tuwien.ac.at

October 23, 2018

Abstract

We present Stamp-it, a new, concurrent, lock-less memory reclamation scheme with amortized, constant-time (thread-count independent) reclamation overhead. Stamp-it has been implemented and proved correct in the C++ memory model using as weak memory-consistency assumptions as possible. We have likewise (re)implemented six other comparable reclamation schemes. We give a detailed performance comparison, showing that Stamp-it performs favorably (sometimes better, at least as good as) than most of these other schemes while being able to reclaim free memory nodes earlier.

1 Introduction

Efficient, dynamic memory management is at the heart of many sequential and parallel algorithms, consisting in the allocation of pieces of memory and the subsequent, *safe* reclamation of these pieces when they are no longer in use. In parallel and concurrent, lock- and wait-free algorithms, the reclamation step is highly non-trivial since more than one *thread* may be referencing and using an allocated piece of memory unbeknownst to other threads: It cannot be given back to the system or thread-local heap before it has been ascertained that no threads will possibly access any data in this memory anymore.

There has been a substantial amount of work on memory reclamation for concurrent algorithms, see, e.g., [1–3,5–10,13,15,16,19,24,26,29]. All of these schemes have their merits and (performance) issues. One drawback shared by them all, except for reference counting schemes, is that they need to scan references from *all threads* in order to reclaim possibly no longer referenced memory pieces. A main motivation of this work is to overcome this bound.

Our contribution is a new lock-less reclamation scheme, called *Stamp-it*, which is compared qualitatively and experimentally to six well-known and, depending on circumstances, well performing current schemes. Reclamation in Stamp-it is done in *amortized constant time* per reclaimed memory block; no references are scanned unless they can be reclaimed. All tested schemes have been (re)implemented in C++; full source code is available at http://github.com/mpoeter/emr. The experimental evaluation is done on four architecturally different systems with large numbers of hardware supported threads, ranging from 48 up to 512. We use standard benchmarks, as well as a new benchmark designed to study memory consumption by reclaimable but unreclaimed memory. On these benchmarks and machines, Stamp-it compares favorably to and in many cases and aspects significantly outperforms the competing schemes.

In the following, a contiguous piece of memory allocated from the system heap for use in a concurrent algorithm and possibly shared between threads is called a *node*. We do not deal with memory management (allocation and deallocation of nodes) here. Nodes may store additional meta-information that is not visible to the application, and we mention where such is required.

We are interested in *general purpose* reclamation schemes that allow eventually reclaimed nodes to be freely reused at a later time, regardless of how and in which data structure the allocated nodes were used. Not all reclamation schemes have this property, e.g., [26, 28] do not allow general reuse of reclaimed nodes, [5, 13] have to be tailored to the application data structure and [7,8] require the data structure to be in a special, normalized form. A general purpose scheme should be *non-intrusive*, requiring no or little changes in the application code. A way of achieving this is to rely on a standard interface as proposed for C++ [25]. A reclamation scheme should be fast, both in use and maintenance of references to shared nodes, as well as in the actual reclamation. It should require little memory overhead, avoid typical performance issues like false sharing and should not prevent applications using commonly found patterns in lock-free programming like borrowing some bits from a pointer. Reclaimability of nodes should be detected fast to reduce unnecessary memory consumption. Robustness against crashes, and bounds on the amount of memory blocked by crashed threads are desirable but not provided by most schemes. We say that a reclamation scheme is *lock-less*, if, provided that sufficient system memory is available, no thread can block the progress of the application. On the other hand, a scheme is reclamation-blocking if a suspended or crashed thread can prevent an unbounded amount of nodes from being reclaimed [15]. With this terminology, Stamp-it is lock-less but reclamation-blocking. Lock-freedom should allow good scalability; wait-freedom would be desirable, but not many schemes actually provide this.

All lock- and wait-free algorithms rely on hardware supported atomic operations. We consider only solutions that use standard atomics like fetch-and-add (FAA) and single-word compare-and-swap (CAS). Solutions requiring non-standard double-word compare-and-swap (as in, e.g., [9]) will either be non-portable or require expensive emulations. We also rule out solutions that have to be tailored to specific data structures like [5, 13], or require hardware or operating system specific features like transactional memory, e.g., [1] or POSIX signals, e.g., [2,6]. The aim was to design a portable, fully C++ standard conform and platform independent implementation. Our implementation is mature beyond a simple proof of concept, and applicable for real-life applications. It works with arbitrary numbers of threads that can be started and stopped arbitrarily.

Based on the above discussion, we have implemented six comparable schemes that make the same assumptions. These are Lock-free Reference Counting (LFRC) [28], Hazard Pointers (HP) [19], Quiescent State-based Reclamation (QSR) [17], Epoch-based Reclamation (ER) [12], New Epoch-based Reclamation (NER) [15], and DEBRA [6]. Interval-based Reclamation (IR) would fit among these, but is too recent to be considered [29]. Hart et al. [15] used a similar selection of schemes in their study, and we wanted to repeat their experiments with our own implementations on different platforms and at a larger scale. Common to all implementations, including Stamp-it, is that we rely on the C++ memory model and can argue that our implementations are correct in this model. We try to use as weak consistency assumptions as possible.

In the remainder of this report we first illustrate how a reclamation scheme for C++ can be used in lock-free algorithms and data structures with the C++ interface that we support [25]. In Section 3 we describe our new reclamation scheme Stamp-it and argue for correctness and amortized complexity. An experimental comparison between Stamp-it and the other six schemes is given in Section 4. More results can be found in the appendix, a previous technical report [22, 23] and in [21] which is the basis for this work.

2 C++ Memory Reclamation Interface

Before presenting the reclamation schemes, we illustrate how memory reclamation can be done in actual applications using an interface proposed for C++ by Robison [25]. This proposal defines an abstract interface and allows many different reclamation schemes to be implemented and used. The interface defines the following fundamental pointer abstractions:

- A marked_ptr allows one or more low-order bits to be borrowed. Many lock-free algorithms rely on such mark tricks, e.g., [4,14,27]. The get method returns the raw pointer (without the mark bits), and the mark method returns the value of the mark bits.
- A concurrent_ptr acts like an atomic marked_ptr, i.e., it supports atomic operations.
- A guard_ptr is an object that can atomically take a snapshot of the value of a concurrent_ptr and guarantee that the target will not be deleted as long as the guard_ptr holds a pointer to it.

It is important to note that only guard_ptr references protect against deletion of a node. In effect, a concurrent_ptr is a "weak" pointer and a guard_ptr is a "shared ownership" pointer.

To obtain a snapshot from a concurrent_ptr and populate a guard_ptr, acquire and acquire_if_equal methods can be used. In wait-free algorithms, acquire may be problematic with some schemes like Hazard Pointers, or even LFRC, because it may have to loop indefinitely. For such cases, acquire_if_equal can be used as it simply stops trying if the value in the pointer does not match an expected value, and reports whether it was successful or not.

Releasing a guard_ptr follows the standard smart pointer interface. For a guard_ptr instance g, the operation g.reset releases ownership and sets g to nullptr; the guard_ptr destructor implicitly calls reset. The reclaim method marks the given node to be reclaimed once it is safe to do so and implicitly resets the guard_ptr.

An implementation of the find function in Harris' list-based set [14] with the improvements proposed by Michael [18] using this interface is shown in Listing 1.

Listing 1: Implementation of list::find

```
template <class Key, class Reclaimer>
    bool list<Key, Reclaimer>::find(Key key,
   concurrent_ptr*& prev, marked_ptr& next,
       guard_ptr& cur, guard_ptr& save) {
    retry:
       next = prev->load();
       save.reset();
           if (!cur.acquire_if_equal(*prev, next))
          goto retry;
if (!cur) return false;
next = cur->next.load();
\begin{array}{c} 12 \\ 13 \\ 14 \\ 15 \\ 16 \\ 17 \\ 18 \\ 19 \\ 20 \\ 21 \\ 22 \\ 23 \\ 24 \\ 25 \\ 26 \\ 27 \\ 28 \end{array}
           if (next.mark() != 0) {
              next = cur->next.load().get();
             cur.reclaim();
              if (prev->load() != cur.get()) goto retry;
             Key ckey = cur->key;
if (ckey >= key) return ckey == key;
             prev = &cur->next;
save = std::move(cur);
       }
```

We first acquire a guard_ptr to the next node and store it in cur, ensuring that we can safely iterate the list. We use the acquire_if_equal method since we already know the expected value. We then check if cur's next pointer has the mark bit set. In that case we try to splice out the

node, and mark cur for reclamation. Otherwise, we continue to iterate the list. We move cur into save, i.e., the previous guard_ptr in save (if any) gets reset and replaced with the value from cur; and cur is reset.

We made a number of small changes and adaptations to Robisons interface proposal [25]. These changes are described in detail in [21]. The most important change is the introduction of the region_guard concept, which is required for reclamation schemes like NER, QSR and Stampit. In these schemes a guard_ptr can only exist inside a critical region, so unless the thread is already inside a critical region the guard_ptr automatically enters one. Entering and leaving critical regions are usually rather expensive operations, and region_guards allow to amortize this overhead. Any guard_ptr instances created inside the scope of a region_guard can simply use the current critical region and save the overhead of entering a new one. In QSR, the region_guards are used to reduce the number of fuzzy barriers.

3 Stamp-it Memory Reclamation

We now introduce our new scheme, Stamp-it. It is an epoch-based scheme conceptually similar to NER and therefore provides many of the same properties. As in ER/NER, the programmer has to define critical regions that are entered and left explicitly. A thread is only allowed to access shared objects inside such regions.

The Stamp-it algorithm maintains thread-local and one global *retire-list* of nodes that can potentially be reclaimed, and relies on an abstract data structure called *Stamp Pool* that efficiently supports the following operations:

- 1. Add an element and assign a stamp to it (push). Stamps have to be strictly increasing, but not necessarily consecutive.
- 2. Remove a specific element, return true if this element was the one with the lowest stamp at that point in time (remove).
- 3. Get the highest stamp assigned to an element so far, i.e., the last stamp that has been assigned to an element.
- 4. Get the lowest stamp of all elements currently in the Stamp Pool.

The algorithm uses the Stamp Pool as follows. Upon entering a critical region the thread adds itself to the pool, and gets a new stamp value. Stamp values are strictly increasing, therefore, defining a total order in which all threads have entered their respective critical region. When a thread retires a node it requests the highest stamp from the Stamp Pool, stores it in the node and appends the node to the end of its local retire-list.

Upon leaving a critical region the thread removes itself from the Stamp Pool. and performs a reclaim operation on the local retire-list. The reclaim operation requests the lowest stamp from the Stamp Pool. It can then safely reclaim all nodes with a smaller stamp value. Since nodes are appended to the end, the elements in the local retire-list are ordered by their stamp values. This makes reclamation very efficient as it always has a runtime linear in the number of nodes that can currently be reclaimed; no time is wasted on nodes that cannot yet be reclaimed. The working of Stamp-it is illustrated in Figure 1.

If the remove operation returns false and the number of nodes in the local retire-list exceeds some threshold, the thread pushes all remaining entries to the global retire-list as an ordered sublist. If remove returns true, i.e., the thread was the last one, it performs a reclaim operation on the global retire-list. The global retire-list is not ordered and therefore does not provide the same runtime guarantees. However, since it is organized as a list of ordered sublists, each

sublist needs to be scanned only up to the node which has a stamp that is larger than or equal to the lowest stamp. Therefore, the resulting total runtime is O(n+m) where n is the total number of reclaimable nodes and m is the number of ordered sublists.

We will argue for the following propositions, and deal with progress properties later:

Proposition 1 Stamp-it is reclamation safe: A node is reclaimed only when it is referenced by no thread.

Assume that some thread retires (want to reclaim) some node n. It fetches the currently highest stamp from the Stamp Pool, stores this stamp value in n and adds n to its local retire-list. Now, n can safely be reclaimed once all threads that were in a critical region at the time n was removed have left their respective critical regions. The lowest stamp in the Stamp Pool is smaller than or equal to the stamp of the last thread, so n's stamp being less than this lowest stamp implies that all threads currently inside critical regions (if any) have entered their respective critical region $after\ n$ was retired, and therefore n can safely be reclaimed (no references).

Proposition 2 Stamp-it reclaims any node in amortized constant time in the number of Stamp Pool operations.

Assume that q threads are inside critical regions, when thread T_1 is leaving its critical region and thus removing itself from the Stamp Pool. All nodes in T_1 's local retire-list can be reclaimed once these q threads have left their critical regions. Thread T_1 can possibly add its local retire-list as sublist to the global retire-list, deferring reclamation of its nodes to the *last* thread to leave its critical region. The global retire-list is traversed only when this last thread removes itself from the Stamp Pool. In the worst case, the q threads remove themselves in the same order as they have entered their critical regions, resulting in q traversals of the global retire-list. We can therefore achieve amortized constant time by only adding local retire-lists that hold more than q nodes (threshold); the sublist might be touched up to q times to reclaim at least q nodes.

Maintaining a dynamic counter for q can cause additional overhead. Alternatively, one could use a less volatile counter as upper bound such as the total number of threads, but then the total number of nodes left in local retire-lists would be quadratic in the number of threads and would therefore increase unnecessary memory consumption. For this reason, we use a static threshold with an empirical value of 20 in our implementation.

Our implementation of the Stamp Pool is built on the ideas of the lock-free doubly-linked list by Sundell and Tsigas [27]. It requires two static dummy nodes, head and tail, which are also used to manage the highest and lowest stamp values; the highest stamp is stored in head and the lowest one in tail.

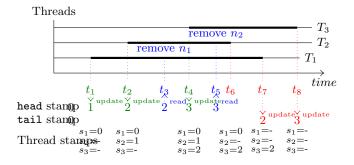


Figure 1: An example sequence of Stamp Pool operations. Thick lines indicate critical regions.

Our implementation differs from [27] in that we only need to push a node from one direction (next to head), and that each node can be removed at any time, independent of its position. Each thread has a single local thread_control_block that acts as a node in the Stamp Pool, i.e., the nodes are "reused" and we therefore have to take care of the ABA problem and consider the possibility that nodes might reappear at different positions. We will from now on refer to the thread_control_blocks in the Stamp Pool as blocks (not to be confused with reclaimable memory nodes).

To insert or delete a block from the Stamp Pool one has to update the respective set of prev and next pointers. These have to be changed consistently, but not necessarily all at once. The solution proposed by Sundell and Tsigas is to treat the doubly-linked list as a singly-linked list with auxiliary information in the prev pointers. Thus, the next pointers always form a consistent singly-linked list, but the prev pointers only give hints for where to find the previous block. We use the same approach, but reversed the directions, i.e., keep the prev list consistent and use the next pointers as auxiliary information, as this is more suited to the use in Stamp-it.

Both pointers, next and prev have to be equipped with a deletion mark (in the least significant bit) to prevent conflicting updates from concurrent insert and delete operations as in Harris' singly-linked list [14].

To avoid the ABA problem, in addition to the delete mark we spare additional 17 bits for a version tag in both pointers. These bits are used to store a tag that gets incremented with every change to the pointer value. There is still a very small chance for an undetected ABA to occur when the version tag wraps around, but in order for this to happen there have to be exactly 2^{17} updates to the pointer between the initial read and the subsequent CAS operation.

3.1 Stamp Pool data structure

The Stamp Pool keeps track which threads have entered a critical region and in which order. It holds two static dummy blocks head and tail; new blocks are inserted right after head. The prev pointers define the direction from head to tail; this direction is always kept consistent. The next pointers define the direction from tail to head, but only act as hints where to find the next block. It is therefore possible that a block, which is already in the prev list, does not occur (yet) in the next list (and the other way round in case of removal).

Each queue block, including head and tail, holds a stamp counter. When a new block is inserted, it loads head's prev pointer and stores it in its local prev, increases head's stamp using an FAA operation, stores the returned value in its local stamp and then performs a CAS on head->prev in order to insert itself into the prev list. This ensures that head always holds the highest stamp and that the stamps in the prev direction are strictly decreasing, i.e., the stamp of the newly added block is greater than all other blocks (except head). The only exception to this is the tail block which should always reflect the stamp value of its immediate predecessor in the prev direction. This way we can easily fetch the lowest stamp value from tail.

Even though the next pointers only act as hints, it is guaranteed that they only point to blocks with a higher stamp value, i.e., the stamps in the next direction are strictly increasing (with the tail block again being an exception). More specifically, the next pointer of some block b can point to:

- head: This can be the case when head is the predecessor of b, or when some other block c has inserted itself between head and b, but did not yet update b's next pointer. Note that when b has already marked its prev pointer, it can no longer be updated by c, so this inconsistency can only be resolved once b is fully removed from both lists.
- a block c which is still in the prev list: This is the "normal case", i.e., usually c is the

predecessor of b; unless b is head, in which case the previously described exception is possible.

• a block c which is removed from the prev list: This is the intermediate state when c has been removed from the prev list, but not yet from the next list. So prev->next points to c, but c is no longer the immediate predecessor of b in the prev direction when starting from head. However, by following c's next pointer (and potentially those of other removed blocks) one can find b's new predecessor.

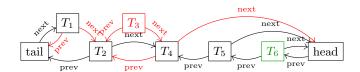


Figure 2: Example of blocks with their links; green blocks are not yet fully inserted, red blocks are removed; red links are marked.

An example is shown in Figure 2. The links of the blocks T_1 , T_3 and T_4 are all marked so they cannot be updated. The block T_3 is already fully removed, i.e., it is not referenced by any prev nor next pointer in the list. The blocks T_1 and T_4 are marked for deletion, but are not yet fully removed; T_1 has been removed from the prev list, but is still in the next list; T_4 is still fully linked. Block T_5 already finished its push operation (not green), but the next pointer of its successor, T_4 , still points to head. This indicates that T_4 's next pointer was already marked, so it could not be updated by T_5 . On the other hand, T_6 is currently in the push operation (green); it has successfully inserted itself in the prev list, but the update of T_5 's next pointer is still pending.

The two lowest bits of the stamp counter are used to embed flags to track a block's state:

- PendingPush: The block is currently being inserted into the prev list.
- NotInList: The block has been completely removed, i.e., is no longer part of neither prev nor next list. This implies that the owning thread is no longer inside a critical region.

The two flags are mutually exclusive, so they cannot both be set at the same time. A block can be in four different states:

- in the process of being inserted into the prev list: The PendingPush flag is set, the NotInList flag and next's delete mark are cleared and prev's delete mark is undetermined.
- in the queue: The delete marks of the prev and next pointer, as well as the NotInList flag are cleared, the PendingPush flag is undetermined. Note that this state only indicates that the block was correctly inserted in the prev direction, but does not state anything about the next direction since we cannot update the next pointer of a block that is already marked.
- in the process of being removed: The delete mark of the prev pointer is set, the mark of next is undetermined and the NotInList and PendingPush flags are cleared.
- fully removed: The NotInList flag and prev's delete mark are set, the PendingPush flag is cleared and next's delete mark is undetermined (because it gets reset in the push operation before the PendingPush flag is set).

3.2 The Basic Steps in Detail

We now show how to support the operations on the doubly-linked list data structure correctly in a lock-free manner. Some code listings are included in the main text, the rest can be found in the appendix. For brevity, we omitted the C++ memory orderings from the code shown here. Full code listings with memory orderings and more detailed explanations are provided in [21].

To insert a block we first set the block's next pointer to head implicitly clearing the DeleteMark of next. Then we perform an FAA on head->stamp, getting the new stamp for the block we are about to insert. This new stamp value is modified to have the PendingPush flag set before it is stored in our block. Then we set the block's prev pointer and attempt a CAS operation to update head->prev with our own block. The implementation of the push method is shown in Listing 4 in the appendix. When the CAS is successful our block has been inserted in the prev list. We can therefore reset the PendingPush flag, and perform a final CAS-loop to update our successor's next pointer.

When a thread leaves a critical region, it removes itself from the queue. The remove operation first marks the prev and next before removing the block from both lists. Then it sets the NotInList flag and checks if this thread was the last one, i.e., if the block's prev pointer points to tail. If that is the case, it tries to update tail's stamp to that of the new "last" thread.

Marking the two pointers signals to other threads that this block is about to be removed, and also prevents the pointers from being updated by CAS operations from threads that did not yet see the mark. In order to remove a block b from the prev list, the thread has to find its predecessor, i.e., the block c with the prev pointer pointing to b, and update c's prev pointer with the value of b's prev pointer. But it can of course happen that c's prev pointer is also marked and can therefore not be updated. In this case we have to find c's predecessor and help remove c before we can continue with the removal of b. By removing c, we get a new predecessor for b. We can then restart the loop and try to remove it again. The same idea is applied when removing a block from the next list.

Since a block b can only be removed from the prev list when its immediate predecessor is not marked, any marked immediate predecessor has to be removed before b can be removed. Therefore, whenever a thread that tries to remove a marked block b encounters another block b which is supposed to come after b in the prev direction (i.e., it was found by following the prev pointers starting from b), where stampb0 stampb0 or stampb0 has the NotInList flag set, and all blocks on the path from b1 to b2 are also marked, we can conclude that b3 has already been removed from both lists.

Since c was encountered after b in the prev direction, it is supposed to have a lower stamp than b; it can only have a larger stamp if it was removed and then reinserted. But since all blocks between c and b are marked, c could not have been removed without first removing all those blocks, including b. The same holds for the case when the NotInList flag is set, as the flag is only set once the block has been fully removed.

The code to remove a block from the prev list is shown in Listing 2. The remove operation keeps track of three different pointers:

prev is a reference to the next *unmarked* block in the prev direction, i.e., the block that we want to set as the new value for our predecessor's prev pointer. We get to this block by following our own prev pointer and the prev pointers of other marked blocks (if any).

next is a reference to some block that precedes our own block in the prev direction. By following this block's prev pointer we should end up at our own block, unless some other thread has removed it already. This way we can efficiently find our immediate predecessor to update its prev pointer.

last is a reference to a helper block that is used to remove potentially marked predecessors of our own block. When this pointer is not null, it should be the immediate predecessor of the next block in the prev direction.

So the order of the blocks in the prev direction should be as follows: last (if it is set), next, our own block b, and prev. Each of these blocks (except b) can potentially be removed and reinserted at any time. For next and last we have to consider this possibility and take appropriate actions. However, when we recognize that prev has been removed or reinserted, we can stop since we know that b must have been removed already as well.

Listing 2: Stamp-it's remove_from_prev_list method

```
bool remove_from_prev_list(
   marked_ptr& prev, marked_ptr b, marked_ptr& next)
 3
      const auto my_stamp = b->stamp.load();
      marked_ptr last = nullptr;
      for (;;) {
        if (,,, ) {
  if (next.get() == prev.get()) {
    next = b->next.load(); return false;
 7
8
9
        auto prev_prev = prev->prev.load();
auto prev_stamp = prev->stamp.load()
10
                               prev->stamp.load();
        if (prev_stamp>my_stamp || prev_stamp&NotInList)
13
           return true;
14
            (prev_prev.mark() & DeleteMark) {
15
16
17
18
19
               (!mark_next(prev, prev_stamp)) return true;
           prev = prev->prev.load(); continue;
         auto next_prev = next->prev.load();
        auto next_prev = next->stamp.load();
if (next_prev != next->prev.load()) continue;
20
21
22
23
24
25
26
27
28
            (next_stamp < my_stamp) {</pre>
           next = b->next.load(); return false;
        if (next_stamp & (NotInList | PendingPush)) {
           if (last.get() != nullptr)
              next = last; last.reset();
             else next = next->next.load():
           continue;
29
30
        if (remove_or_skip_marked_block(
31
                next, last, next_prev, next_stamp))
32
33
            (next_prev.get() != b.get()) {
34
           move_next(next_prev, next, last); continue;
35
36
        if (next->prev.CAS(next_prev, prev))return false;
```

The remove_from_prev_list operation essentially consists of a large loop that keeps track of the three mentioned blocks, while trying to find the direct predecessor of b. Once that predecessor is found, we can try to update its prev pointer in order to remove b. There are several conditions that lead to the termination of this loop. In some of these cases we can conclude that b is already removed from both lists, in other cases we know that b has been removed from the prev list, but we still need to ensure that it is also removed from the next list.

In line 7, if prev and next point to the same block, b must have been removed from the prev list already. If in line 12 the prev's stamp is greater than b's stamp or has the NotInList flag set, we can conclude that prev must have been removed (together with b).

If prev's prev pointer is marked (line 14), we try to help setting the delete mark on prev's next. The mark_next operation simply performs a CAS loop trying to set the delete mark on the next pointer as long as that block's stamp matches the given stamp value. When we detect that the stamp has changed, we can conclude that prev must have been removed (together with b).

If next's stamp is less than b's stamp (line 21), we can conclude that b must have been removed from the prev list. Otherwise, if next's stamp has the NotInList or PendingPush flag set (line 24) we cannot use this block since it might not be part of the prev list. For the NotInList flag this is clear, but for the PendingPush flag this is more subtle: The flag signals that the block is currently

getting inserted into the prev list, but with the information we have available at this time it is impossible to tell whether this has already happened or not.

In line 30 we call the helper function remove_or_skip_marked_block (shown in Listing 8 in the appendix). This method checks whether the *next* block is marked, and if so, tries to remove it provided we have a valid *last* pointer (recall that *last* is supposed to be the predecessor of *next*). In case we have no *last* pointer, we move *next* to the next block in the next direction.

If next's prev pointer does not match b (line 33) it is not b's predecessor. In that case we call move_next (shown in Listing 3) and restart the loop.

The move_next method tries to move next to the following block in the prev direction, while keeping the old value of next in last. There is a special case that needs to be handled. It could happen that the next block in the prev direction has successfully inserted itself into the list, but still has the PendingPush flag set, i.e., it did not yet finish its push operation. We previously checked that next's stamp does not have the PendingPush flag set, because otherwise we would have to dismiss the block as we could not determine whether it is already inserted. Now we can conclude that it is in fact part of the prev list, so we help resetting the PendingPush flag. This is necessary to ensure lock-freedom, as otherwise we would iterate infinitely because in the next iteration the previously mentioned check would fail again, resulting in next being moved back in the next direction again.

Listing 3: Stamp-it's move_next method

When we arrive at line 37 we have found b's predecessor and can attempt a CAS on next's prev with our current prev to remove b from the prev list.

In case we could not conclude that b is already removed from both lists, we still have to remove it from the next list; this is done in the remove_from_next_list operation (shown in Listing 6 in the appendix), which is quite similar to remove_from_prev_list. It also keeps track of the same three pointers, where the initial values for prev and next are those that were returned by remove_from_prev_list. This allows us to continue from where we left, reducing the amount of work to find the blocks we need to update in many cases.

In this method we have to set next to the last unmarked block with a stamp greater than b's stamp, and prev to the first unmarked block with a stamp less or equal to b's stamp (both in the prev direction), i.e., the two blocks that would be the predecessor and successor of b if b would still be part of the prev list. This entails that next's prev pointer must reference prev. Once we have found these blocks, we can attempt a CAS to update prev's next pointer in order to finish removal of b from the next list. If the CAS succeeds, we have successfully removed b. However, we still have to make sure that the prev block has not been marked in the meantime. If this is the case, we have to continue and help remove prev from both lists in order to maintain the previously described condition, which lets us conclude that a block has been fully removed if we recognized that the successor block has been fully removed.

The update_tail_stamp method (shown in Listing 9 in the appendix) tries to find the new predecessor of tail in the prev direction, read its stamp and update tail's stamp accordingly. Unfortunately, finding this predecessor is not as simple as taking tail's next pointer, since it

could point to head (due to the predecessor not having finished its push operation) or to a block that could have been removed and potentially reinserted at the time we read its stamp. Of course we could detect such cases and try to find the actual predecessor, but we do not want to waste too much time for this. Instead, if we cannot immediately identify the new predecessor we simply use the "next best guess" for the new stamp, which is our own block's stamp plus a stamp-increment (recall that stamps are strictly increasing).

Finally, we perform a simple CAS-loop, trying to update tail's stamp as long as the new value we want to write is greater than the value we are trying to replace.

3.3 Correctness

We can now argue for the following proposition:

Proposition 3 Stamp-it is lock-less, that is, all methods used for entering and leaving critical regions (push, remove), and all helper functions are lock-free. The expected average runtime of the operations is constant, if threads do not conflict.

Stamp-it is, however, reclamation-blocking in the sense that a stalled (or crashed) thread can prevent an unbounded number of nodes for being reclaimed [15].

The push operation (see Listing 4 in the appendix) is lock-free. The first loop performs a CAS operation in order to insert the block into the queue. In case the CAS succeeds, we break out of the loop, otherwise we just restart the loop. The CAS can only fail if some other thread interfered—either by inserting or removing some block. But in this case some other thread must have made progress. The same argument applies to update_tail_stamp and mark_next. Both methods contain loops that perform CAS operations, and a failure of these operations can only be caused by progress in some other thread.

The remove_from_prev_list (see Listing 2) and remove_from_next_list (see Listing 6 in the appendix) operations are a bit more complex. We argue for remove_from_prev_list, and since both operations are similar, the same will hold for remove_from_next_list. As mentioned, both keep track of a prev and a next pointer. In each iteration we perform one of the following changes in case we have to restart the loop:

- move prev along the prev direction (in case prev is marked)
- move next along the prev direction (in case next is not prev's predecessor)
- remove next from the prev list (in case next is marked and we have a valid last pointer)
- move next along the next direction (in case next is marked and we have no last pointer, or next has the NotInList Or PendingPush flag set)
- nothing (in case the CAS to remove b failed)

The block b divides the prev list into two sublists: The sublist from head to b, and the sublist from b to tail. next points to a block in the first sublist and prev points to a block in the second sublist. New blocks are inserted at the beginning of the prev list (right after head), i.e., they become part of the first sublist. So the number of times we can move next in the prev direction before we reach b is bounded by the number of entries in the first sublist, and the number of times we can move prev in the prev direction before we reach tail is bounded by the number of entries in the second sublist.

The case where we have to move *next* back in the next direction because it is marked and we have no valid *last* can be resolved by following *next*'s next pointer and from there move again

along the prev pointer, while maintaining *last*. So the next time we encounter the same marked block, we will be able to remove it as we now have a valid *last* pointer. In the worst case scenario we have to move along the next direction until *next* points to head, from where we can then start to move *next* along the prev direction, while removing marked blocks (if any). The case where *next* has the PendingPush flag set can be resolved in the same way.

This leaves us with the cases where *next* has the NotInList flag set or the CAS operation to remove b fails. But both cases can only occur when another thread changed the data structure in a way that it is no longer consistent with the thread's view. So unless some other thread interferes, for both methods, remove_from_prev_list and remove_from_next_list, it is guaranteed that at any time a thread is able to finish the method in a bounded number of steps.

Unfortunately, the block pointed to by *next* can be removed and reinserted at any time. Obviously, this destroys the bounds as with every reinsertion the block is put back right at the beginning of the prev list. However, this implies that the owning thread of this reinserted block has been able to finish its remove and subsequent push operation, i.e., that it has made progress. Thus, the requirements for lock-freedom are fulfilled as it is guaranteed that at any time at least one thread makes progress: If there is no conflict with another thread, we can finish the operation in a bounded number of steps; otherwise, the interfering thread was able to make progress.

The push and remove operations are only lock-free and therefore do not provide an upper bound on the number of iterations. In practice however, the average runtime is expected to be small (constant). We have verified this experimentally, but have to refer to [21, 22] for the results due to the limited space.

Finally, we claim the following.

Proposition 4 The implementation is correct under the C++ memory model.

All atomic operations are relaxed as far as possible without sacrificing correctness with the appropriate C++ memory model annotations. Due to the limited space we cannot show full code with memory orderings, which we use to carefully argue that required happens-before relationships hold as needed. However, it is not possible to follow the correctness arguments on the basis of the C++ memory model's semantics without the corresponding code; we have to refer to [21] for full listings and correctness arguments (also for the C++ implementations of the other schemes).

4 Experimental Comparison

We have compared Stamp-it with six other currently considered reclamation schemes: Lock-free Reference Counting (LFRC) [28], Hazard Pointers (HPR) [19], Epoch Based Reclamation (ER) [12], New Epoch Based Reclamation (NER) [15], Quiescent State Based Reclamation (QSR) [17] and DEBRA [6]. All schemes have been implemented in C++ using the adapted interface described in Section 2 and the C++ memory model. All implementations are tuned by relaxing the atomic operations as far as possible. Correctness arguments based on the memory model's semantics for all implementations are provided in [21]. The full source code is available at http://github.com/mpoeter/emr. All results and scripts are available at http://github.com/mpoeter/emr-benchmarks.

The tests are set up similarly to Hart et al. [15] and we also repeat most of those analyses. We can show here only a subset of our results, more results can be found in [21, 22] and in the appendix.

4.1 Benchmarks

We tested the reclamation schemes on a (1) queue, a (2) linked-list and a (3) hash-map. The queue is based on Michael and Scott's design [20], the linked-list and hash-map on Michael's improved version [18] of Harris' list-based set [14]. For the List benchmark the key range is twice the initial list size. The probabilities of inserting and removing nodes are equal, keeping the size of the list and queue data structures roughly unchanged throughout a given run. The List benchmark has a workload parameter that determines the fraction of updates (remove/insert) of the total number of operations. A workload of 0% corresponds to a search-only use case, while a workload of 100% corresponds to an update-only use case.

Our experiments are throughput oriented in the following sense. The main thread spawns p child threads and starts a timer. Every child thread performs operations on the data structure under scrutiny until the timer expires. The parent thread calculates the average execution time per operation by summing up the runtimes of the child threads and their number of performed operations. Each benchmark was performed with 30 trials with eight seconds runtime per trial. Most of the benchmarks focus on performance, and calculate the average runtime per single operation for each trial. Each thread calculates its average operation runtime by dividing its active, overall runtime by the total number of operations it performed. The total average runtime per operation is then calculated as the average of these per-thread runtime values.

The Queue and List benchmarks are synthetic micro-benchmarks, exactly as by Hart et al. [15]. The HashMap benchmark is intended to highlight other properties of the reclamation schemes. It mimics the calculation in a complex simulation where partial results are stored in a hash-map for later reuse. These partial results are relatively large, so in order to limit the total memory usage the number of entries in the hash-map is kept below some threshold by evicting old entries using a simple FIFO policy. The resulting benchmark has the following properties:

- there is no upper bound on the number of nodes that are *intentionally* blocked from reclamation.
- the average lifetime of each guard_ptr is relatively long.
- the memory footprint of each node is significant.

Since there is no upper bound on the number of nodes that need to be available for a thread, we have to use the extended hazard pointer scheme that supports a dynamic number of hazard pointers as explained by Michael [19]. The number of buckets in the hash-map is 2048 and the maximum number of entries is 10000. There are 30000 possible partial results and every thread has to calculate or reuse 1000 partial results per "simulation". The size of a partial result is 1024 bytes.

It is important to note that all 30 trials were performed one after the other within the same process. This is especially important in case of the HashMap benchmark as the hashmap is retained over the whole runtime. This means that a result calculated in the first trial can be found and reused in a subsequent trial. For this reason, performance will be worse at the beginning, while the hash-map is in the "warm up phase", but will improve over time when it becomes filled and more items can be reused. But also in the other benchmarks, it is possible that previous trials have impact on later ones, e.g., due to an already initialized memory manager. It was a deliberate design decision to run all trials in the same process as this might more closely reflect a real world situation.

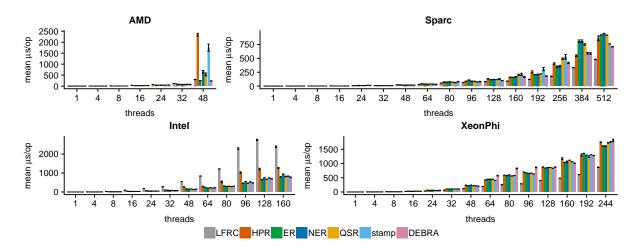


Figure 3: Queue benchmark with varying number of threads.

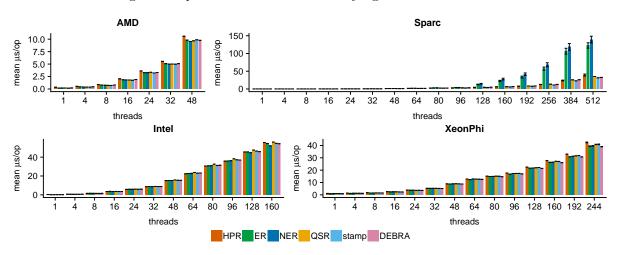


Figure 4: List benchmark with 10 elements, a workload of 20% and varying number of threads (without LFRC).

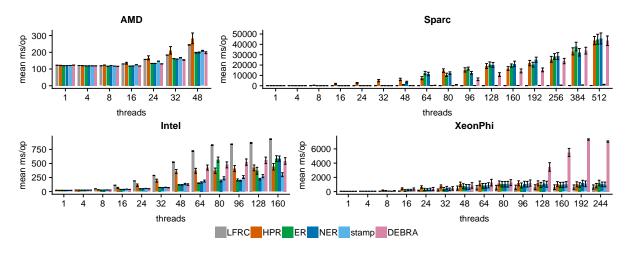


Figure 5: HashMap benchmark with varying number of threads.

4.2 Environment

We executed our tests on four machines with different (micro)architectures. Their respective characteristics are shown in Table 1. These machines all have a relatively large number of cores and hardware supported threads, allowing us to run our experiments at a scale not found in most prior studies. We did not experiment with oversubscribed cores.

Table 1: The four machines used in the experimental evaluation

	CPUs	4x AMD Opteron(tm) Processor 6168
	Frequency	\max . 1.90GHz
\mathbf{AMD}	Cores/CPU	12
	SMT	_
	Hardware Threads	48
	Memory	128GB
	OS	Linux 4.7.0-1-amd64 #1 SMP
		Debian 4.7.6-1 (2016-10-07) x86_64 GNU/Linux
	Compiler	gcc version 6.3.0 20170205 (Debian 6.3.0-6)
	CPUs	8x Intel(R) Xeon(R) CPU E7- 8850
	Frequency	$\max. 2.00 GHz$
\mathbf{Intel}	Cores/CPU	10
	SMT	2x
	Hardware Threads	160
	Memory	1TB
	OS	Linux $4.7.0$ -1-amd $64 \# 1 \text{ SMP}$
		Debian 4.7.6-1 (2016-10-07) x86_64 GNU/Linux
	Compiler	icpc version 17.0.1 (gcc version 6.0.0 compatibility)
	CPUs	1x Intel(R) Xeon Phi(TM) coprocessor x100 family
	Frequency	\max . 1.33 GHz
XeonPhi	Cores/CPU	61
	SMT	4x
	Hardware Threads	244
	Memory	16GB
	OS	Linux 2.6.38.8+mpss3.8.1 #1 SMP
		Thu Jan 12 16:10:30 EST 2017 k1om GNU/Linux
	Compiler	icpc version 17.0.1 (gcc version 5.1.1 compatibility)
	CPUs	4x SPARC-T5-4
	Frequency	\max . $3.60GHz$
SPARC	Cores/CPU	16
	SMT	8x
	SIVII	
	Hardware Threads	512
	·-	
	Hardware Threads	512

We used the jemalloc [11] memory manager on all systems. The main reason being that on Solaris the libc memory manager uses a global lock. For comparison we also ran the

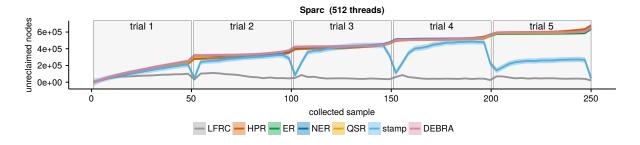


Figure 6: Number of unreclaimed nodes over time in the HashMap benchmark on Sparc. The x-axis is the current sample.

experiments with the standard libc memory manager on all systems except SPARC; these results are shown in Appendix A.3.

ER/NER try to advance the epoch every 100 critical region entry. DEBRA checks the next thread every 20 critical region entries. In the List and Queue benchmarks, a region_guard spans 100 benchmark operations, so this is the size of the critical region for QSR, NER and Stamp-it. QSR executes a fuzzy barrier when it exits the critical region. In HPR, a local retire-list is scanned once its threshold is exceeded; the threshold is $100 + \sum_{i=0}^{p} K_i * 2$ where p is the number of threads and K_i is the number of hazard pointers for the thread with index i.

4.3 Scalability with Threads

We first study the effect of increasing the number of threads that share a single instance of some data structure.

Figure 3 shows the performance of the reclamation schemes in the Queue benchmark. Surprisingly, LFRC performs by far the best on Sparc and on XeonPhi, but is by far the worst on Intel. On AMD, HPR has a huge performance drop when running with the maximum number of threads. A similar effect can be seen by the other schemes as well, especially NER and Stamp-it, but less significant. Apart from these exceptions, all schemes seem to scale largely equally well in this scenario.

For the results of the List benchmark in Figure 4, LFRC has been excluded because it performs exceedingly poor in this scenario. On AMD, Intel and XeonPhi, all schemes are more on less on par, but on Sparc ER and NER show a significant degradation when the number of threads grows beyond 128. We did not investigate the reasons for this in more detail.

Finally, the results for the HashMap benchmark are shown in Figure 5. QSR has been excluded because it scales very poorly on all architectures in this update-heavy scenario. On AMD, ER, NER, Stamp-it and DEBRA scale almost perfectly, while LFRC's and HPR's performance starts to degrade once the number of threads grows beyond 16. On Intel, LFRC scales very poorly while all other schemes scale more or less equally well, but not as well as on AMD. On XeonPhi on the other hand, LFRC scales best while HPR's performance starts degrading with more than 16 threads, but it again improves with more than 128 threads. The other schemes continuously loose performance when the number of threads grows from 16 to \sim 80, but then stays more or less the same. DEBRA's performance drastically breaks down with more than 128 threads.

The biggest surprise is the result on Sparc. Here, the performance of HPR, ER, NER and DEBRA degrades dramatically, while LFRC and Stamp-it scale almost perfectly. With 512 threads the performance difference between LFRC/Stamp-it and the other schemes is a factor of \sim 4000. The reason for this will become clear when we look at the results of the reclamation

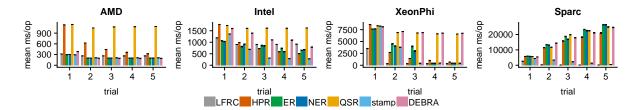


Figure 7: Development of runtime over time in the HashMap benchmark.

efficiency analysis in the next section.

4.4 Reclamation Efficiency

This analysis focuses on how efficiently (fast) the various schemes actually reclaim retired nodes. An increased reclamation efficiency can drastically reduce memory pressure, which in turn can have a significant impact on the overall application performance. Nonetheless, this aspect is usually disregarded in analyses of concurrent reclamation schemes.

To measure reclamation efficiency we use thread-local performance counters that track the number of allocated and reclaimed nodes. By calculating the differences, we get the number of unreclaimed nodes, which is our measurement for efficiency; a smaller number of unreclaimed nodes means that the reclamation scheme works more efficiently.

The plots in this analysis show the development of the number of unreclaimed nodes over time. Each configuration is run with five trials, each with a runtime of eight seconds. During each trial a total of 50 samples are collected. Since the benchmarks are randomized each configuration with the five trials is run 20 times to account for any fluctuation in the measured samples. The plots show the smoothed conditional means of the measured samples of those 20 runs over the number of samples recorded during each run.

For reclamation efficiency, reference counting is the "gold standard". In contrast to all other schemes there is no delay: A node is reclaimed immediately when the last thread drops its reference to that node. So in all the plots, LFRC can be seen as the baseline against which all other schemes have to be measured. One has to keep in mind, though, that LFRC is not a general reclamation scheme, since the reclaimed nodes cannot be returned to the memory manager, but are stored in a global free-list.

Figure 6 shows the results for the HashMap benchmark on Sparc. The number of unreclaimed nodes for HPR, ER, NER, QSR and DEBRA is constantly increasing and does not even go down at the end of the trials when all threads are stopped. When a thread terminates, all schemes add the remaining nodes to a global list. But who is responsible to reclaim them, and when? In Stamp-it the responsibility is transferred to the "last" thread. Other schemes do not have a notion of a "last" thread, so the global retire-list is checked by each thread when it performs reclamation on its local retire-list. When a thread tries to reclaim nodes from the global list it steals the whole list, reclaims all reclaimable nodes and then re-adds the remaining nodes to the global list. This leads to a race during the end of a trial; whoever steals the list might not be able to reclaim all nodes yet, but when the remaining nodes are re-added to the global list, there might be no threads left. Stamp-it mitigates this race as only the last thread reclaims the global list. In addition, we can easily check whether the global stamp has changed since reclamation has started, so we can restart reclamation with the new stamp value. Obviously, the effects of this race get more pronounced the more threads are involved.

The failure to efficiently reclaim nodes increases memory pressure, which has a direct impact on the runtime. Figure 7 shows the development of the runtime over the five trials. On Sparc we can see that the runtime of HPR, ER, NER, QSR and DEBRA is increasing with each trial, while with LFRC and Stamp-it it is decreasing. On the other architectures runtime is decreasing for all schemes except QSR. This would be the expected behavior since more results can be reused once the hash-map has been filled.

The results for the other benchmarks and machines can be found in the appendix.

5 Concluding Remarks

This paper introduced *Stamp-it*, a new, general purpose memory reclamation scheme with attractive features. To the best of our knowledge, this is the first non-reference counting based scheme that does not have to scan all other threads to determine reclaimability of a node. We presented a large scale experimental study, comparing the performance of Stamp-it against six other reclamation schemes on four different architectures in various scenarios. Our empirical results show that Stamp-it matches or outperforms the other analyzed reclamation schemes in almost all cases. All of the analyzed schemes are implemented in portable, standard conform C++, based on the standardized interface proposed by Robison [25]; the full source code is available at https://github.com/mpoeter/emr.

For future work it would be interesting to look for other data structures that could replace the doubly-linked list, i.e., data structures that have less overhead while providing all the required properties. In this context we might also try to relax some of these properties (e.g., use a partial order instead of a strict order for thread entries) in order to reduce contention on the data structure.

References

- [1] D. Alistarh, P. Eugster, M. Herlihy, A. Matveev, and N. Shavit. Stacktrack: An automated transactional approach to concurrent memory reclamation. In *Proceedings of the 9th ACM European Conference on Computer Systems (EuroSys)*, pages 25:1–25:14. ACM, 2014.
- [2] D. Alistarh, W. M. Leiserson, A. Matveev, and N. Shavit. Threadscan: Automatic and scalable memory reclamation. In *Proceedings of the 27th ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 123–132. ACM, 2015.
- [3] O. Balmau, R. Guerraoui, M. Herlihy, and I. Zablotchi. Fast and robust memory reclamation for concurrent data structures. In *Proceedings of the 28th ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 349–359. ACM, 2016.
- [4] H. Boehm. An almost non-blocking stack. In *Proceedings of the 23rd Annual ACM Symposium on Principles of Distributed Computing (PODC)*, pages 40–49. ACM, 2004.
- [5] A. Braginsky, A. Kogan, and E. Petrank. Drop the anchor: Lightweight memory management for non-blocking data structures. In *Proceedings of the 25th Annual ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 33–42. ACM, 2013.
- [6] T. A. Brown. Reclaiming memory for lock-free data structures: There has to be a better way. In Proceedings of the 2015 ACM Symposium on Principles of Distributed Computing (PODC), pages 261–270. ACM, 2015.
- [7] N. Cohen and E. Petrank. Automatic memory reclamation for lock-free data structures. In *Proceedings of the 2015 ACM SIGPLAN International Conference on Object-Oriented*

- Programming, Systems, Languages, and Applications (OOPSLA), pages 260–279. ACM, 2015.
- [8] N. Cohen and E. Petrank. Efficient memory management for lock-free data structures with optimistic access. In *Proceedings of the 27th ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 254–263. ACM, 2015.
- [9] D. L. Detlefs, P. A. Martin, M. Moir, and G. L. S. Jr. Lock-free reference counting. In Proceedings of the 20th Annual ACM Symposium on Principles of Distributed Computing (PODC), pages 190–199. ACM, 2001.
- [10] D. Dice, M. Herlihy, and A. Kogan. Fast non-intrusive memory reclamation for highly-concurrent data structures. In *Proceedings of the 2016 ACM SIGPLAN International Symposium on Memory Management (ISMM)*, pages 36–45, 2016.
- [11] J. Evans. A scalable concurrent malloc(3) implementation for freebsd, 2006.
- [12] K. Fraser. *Practical lock-freedom*. PhD thesis, University of Cambridge Computer Laboratory, 2004.
- [13] A. Gidenstam, M. Papatriantafilou, H. Sundell, and P. Tsigas. Efficient and reliable lock-free memory reclamation based on reference counting. *IEEE Transactions on Parallel and Distributed Systems*, 20(8):1173–1187, 2009.
- [14] T. L. Harris. A pragmatic implementation of non-blocking linked-lists. In Proceedings of the 15th International Conference on Distributed Computing (DISC), volume 2180 of Lecture Notes in Computer Science, pages 300-314. Springer-Verlag, 2001.
- [15] T. E. Hart, P. E. McKenney, A. D. Brown, and J. Walpole. Performance of memory reclamation for lockless synchronization. *Journal of Parallel and Distributed Computing*, 67(12):1270–1285, 2007.
- [16] M. Herlihy, V. Luchangco, and M. Moir. The repeat offender problem: A mechanism for supporting dynamic-sized, lock-free data structures. In *Proceedings of the 16th Inter*national Conference on Distributed Computing (DISC), volume 2508 of Lecture Notes in Computer Science, pages 339–353. Springer-Verlag, 2002.
- [17] P. E. Mckenney and J. D. Slingwine. Read-copy update: Using execution history to solve concurrency problems. In *Proceedings of the 1998 International Conference on Parallel and Distributed Computing and Systems*, pages 509–518, Las Vegas, NV, USA, 1998.
- [18] M. M. Michael. High performance dynamic lock-free hash tables and list-based sets. In *Proceedings of the 14th Annual ACM Symposium on Parallel Algorithms and Architectures* (SPAA), pages 73–82. ACM, 2002.
- [19] M. M. Michael. Hazard pointers: Safe memory reclamation for lock-free objects. *IEEE Transactions on Parallel and Distributed Systems*, 15(6):491–504, 2004.
- [20] M. M. Michael and M. L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In *Proceedings of the 15th Annual ACM Symposium on Principles of Distributed Computing (PODC)*, pages 267–275. ACM, 1996.
- [21] M. Pöter. Effective memory reclamation for lock-free data structures in C++. Master's thesis, TU Wien, Vienna, Austria, 2018.

- [22] M. Pöter and J. L. Träff. A new and five older concurrent memory reclamation schemes in comparison (stamp-it), 2017. arXiv:1712.06134.
- [23] M. Pöter and J. L. Träff. Poster: Stamp-it, amortized constant-time memory reclamation in comparison to five other schemes. In 23rd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP), pages 413–414. ACM, 2018.
- [24] P. Ramalhete and A. Correia. Brief announcement: Hazard eras non-blocking memory reclamation. In *Proceedings of the 29th ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 367–369. ACM, 2017.
- [25] A. D. Robison. Policy-based design for safe destruction in concurrent containers. C++ standards committee paper, http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2013/n3712.pdf, 2013.
- [26] H. Sundell. Wait-free reference counting and memory management. In *Proceedings of the* 19th IEEE International Parallel and Distributed Processing Symposium (IPDPS), pages 24b–24b. IEEE, 2005.
- [27] H. Sundell and P. Tsigas. Lock-free deques and doubly linked lists. *Journal of Parallel and Distributed Computing*, 68(7):1008–1020, 2008.
- [28] J. D. Valois. *Lock-Free Data Structures*. PhD thesis, Rensselaer Polytechnic Institute, 1995.
- [29] H. Wen, J. Izraelevitz, W. Cai, H. A. Beadle, and M. L. Scott. Interval-based memory reclamation. In 23rd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP), pages 1–13. ACM, 2018.

A Additional Code and Results

In this appendix, we expand on some of the (mostly implementation) details that were only briefly discussed in the main text. We also provide additional benchmark results on the reclamation efficiency that were omitted from the main text due to space limitations.

A.1 Code for important methods

The implementation of the push method is shown in Listing 4.

Listing 4: The push method

```
push(thread_control_block* block)
        block->next.store(head);
         marked_ptr my_prev, head_prev = head->prev.load();
         size_t stamp;
         for (;;) {
            marked_ptr head_prev2 = head->prev.load();
if (head_prev != head_prev2) {
  head_prev = head_prev2; continue;
9
10
           stamp = head->stamp.fetch_add(StampInc);
block->stamp.store(stamp-(StampInc-PendingPush));
if (head->prev.load() != head_prev) continue;
my_prev = head_prev;
12
            block->prev.store(my_prev);
15
16
            if (head->prev.CAS(head_prev, block)) break;
        block -> stamp.store(stamp);
        auto link = my_prev->next.load();
for (;;) {
  if (link.get() == block ||
18
19
20
21
22
                  link.mark() & DeleteMark ||
block->prev.load() != my_prev) ||
my_prev->next.CAS(link, block))
23
24
               break;
        }
```

The implementation of the remove operation is shown in Listing 5.

Listing 5: The remove method

```
bool remove(marked_ptr block)
{
    marked_ptr prev = set_mark_flag(block->prev);
    marked_ptr next = set_mark_flag(block->next);
    bool fully_removed = remove_from_prev_list(prev, block, next);
    if (!fully_removed)
        remove_from_next_list(prev, block, next);
    auto stamp = block->stamp.load();
    block->stamp.store(stamp + NotInList);
    bool was_last = block->prev.load().get() == tail;
    if (was_last) update_tail_stamp(stamp + StampInc);
    return wasTail;
}
```

Listing 6 shows the implementation of the remove_from_next_list helper function.

Listing 6: The remove_from_next_list method

```
void remove_from_next_list(
      marked_ptr prev, marked_ptr removed, marked_ptr next)
      const auto my_stamp = removed->stamp.load();
      marked_ptr last = nullptr;
for (;;) {
         auto next_prev = next->prev.load()
         auto next_stamp = next->stamp.load();
if (next_prev != next->prev.load()) continue;
         if (next_stamp & (NotInList | PendingPush)) {
   if (last.get() != nullptr) {
      next = last; last.reset();
}
10
11
12
13
14
            else next = next->next.load();
15
            continue;
16
         auto prev_next = prev->next.load();
auto prev_stamp = prev->stamp.load();
         if (prev_stamp > my_stamp ||
```

```
20
21
22
               prev_stamp & NotInList)
             return;
         if (prev_next.mark() & DeleteMark) {
23
24
25
26
27
28
29
30
            prev = prev->prev.load(); continue;
         }
          if (next.get() == prev.get()) return;
         if (remove_or_skip_marked_block(
                 next, last, next_prev, next_stamp))
             continue;
         if (next_prev.get() != prev.get()) {
  move_next(next_prev, next, last); continue;
31
32
33
         if (next_stamp <= my_stamp ||
    prev_next.get() == next.get())</pre>
34
35
36
            return;
         if (next->prev.load() == next_prev &&
               prev->next.CAS(prev_next, next) && (next->next.load().mark() & DeleteMark) == 0)
38
39
      }
    }
```

Listing 7 shows the implementation of the move_next helper function (used in remove_from_prev_list and remove_from_next_list).

Listing 7: The mark_next method

The implementation of the remove_or_skip_marked_block helper function is shown in Listing 8.

Listing 8: The remove_or_skip_marked_block method

Listing 9 shows the implementation of the update_tail_stamp method.

Listing 9: The update_tail_stamp method

```
20
21 auto tail_stamp = tail->stamp.load();
22 while (tail_stamp < stamp) {
23    if (tail->stamp.compare_exchange_weak(tail_stamp, stamp))
24    break;
25  }
26 }
```

A.2 Reclamation efficiency

This section contains additional results that could not be included in the main text. The results are shown in Figures 8, 9, 10, and 11. What can be seen in all scenarios is that HPR's efficiency is inversely proportional to the number of threads. This is due to the threshold for the local retire-list being calculated to achieve amortized constant time. But this causes the number of unreclaimed nodes in the local retire-lists to be quadratic in the number of threads. Even for the Queue benchmark (Figure 8) and List benchmarks (Figures 9 and 10) show this behavior, even though the number of hazard pointers per thread is constant in these scenarios. In the HashMap benchmark (Figure 11) a dynamic number of hazard pointers is used, which makes the situation even worse.

The situation is similar for DEBRA. In order to advance the global epoch, DEDRA does not check all p threads at once, but only checks a single thread on each critical region entry, thus distributing the costs over p critical regions. But obviously with a large number of this significantly delays the update of the global epoch, resulting in poor reclamation efficiency.

In the Queue benchmark (see Figure 8) on AMD QSR and Stamp-it perform relatively bad. The results for the other architectures are dominated by the bad results of DERBA. On Sparc, HPR performs similarly bad as DEBRA.

In the List benchmark (see Figures 9 and 10) DEBRA and HPR perform significantly worse than the other schemes on all architectures. On Sparc, HPR performs by far the worst.

In the HashMap benchmark (Figure 11) we can see that QSR basically fails completely to reliably reclaim nodes on all the architectures. The number of nodes is constantly increasing and does not even go down at the end of the trials when all threads are stopped. This is also the reason why QSR showed such bad performance in the previous analysis in Section 4.3. DEBRA performs quite good on AMD, but very poor on the other architectures.

For HPR we can also see a consistent increase in the number of unreclaimed nodes over time, even though this number sharply drops right at the beginning of a new trial, but also increases again very rapidly. The only exception is Sparc, where no such drop occurs and the number of nodes is increasing all the time. The other schemes all perform relatively good on all architectures; the exception again being Sparc. On Sparc HPR, ER, NER QSR and DEBRA are all performing equally bad. The number of unreclaimed nodes is constantly increasing and does not even go down at the end of the trials when all threads are stopped. This effect is probably caused by the fact that in these schemes every thread is responsible for reclaiming its own retired nodes. In Stamp-it we know if there is some other thread lagging behind, so we can add nodes to a global list and let that thread take responsibility for reclaiming them. This allows Stamp-it to more reliably reclaim nodes, especially at the end of each trial.

A.3 Results for libc

This section contains the same results as shown in Section 4 and Appendix A.2, but using the standard libc memory manager on AMD, Intel and XeonPhi. On Sparc we still used jemalloc since the libc memory manager on Solaris uses a global lock.

The results do not show significant differences. The overall performance is somewhat lower compared to the jemalloc results, especially on Intel, but the distribution of the measured

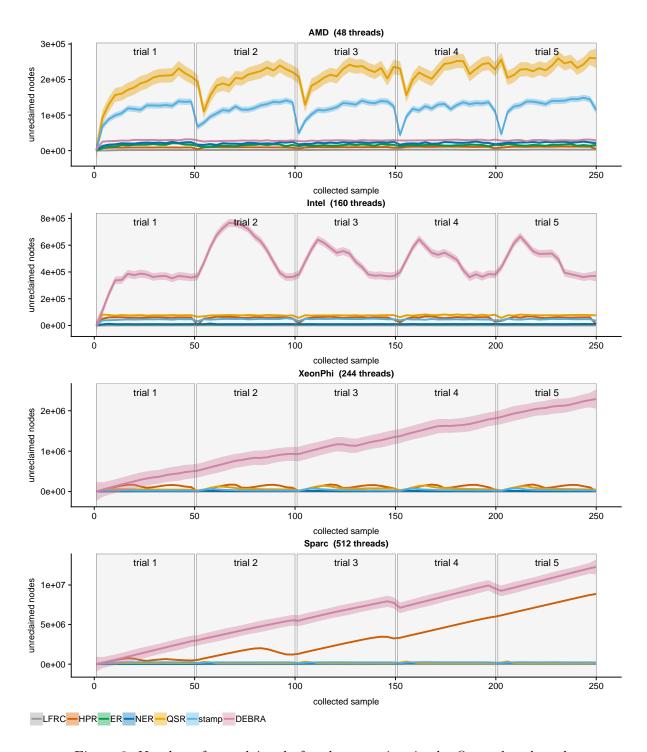


Figure 8: Number of unreclaimed of nodes over time in the Queue benchmark.

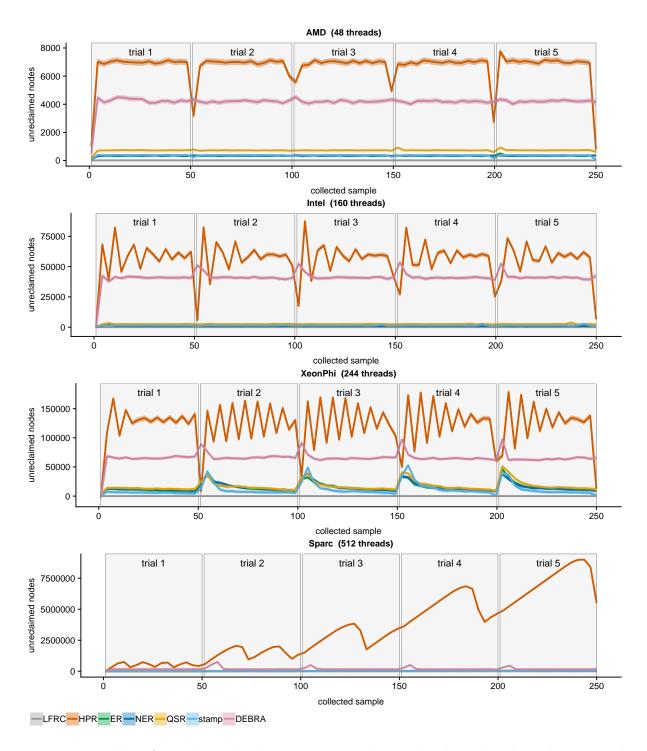


Figure 9: Number of unreclaimed nodes over time in the List benchmark with 10 elements and a workload of 20%.

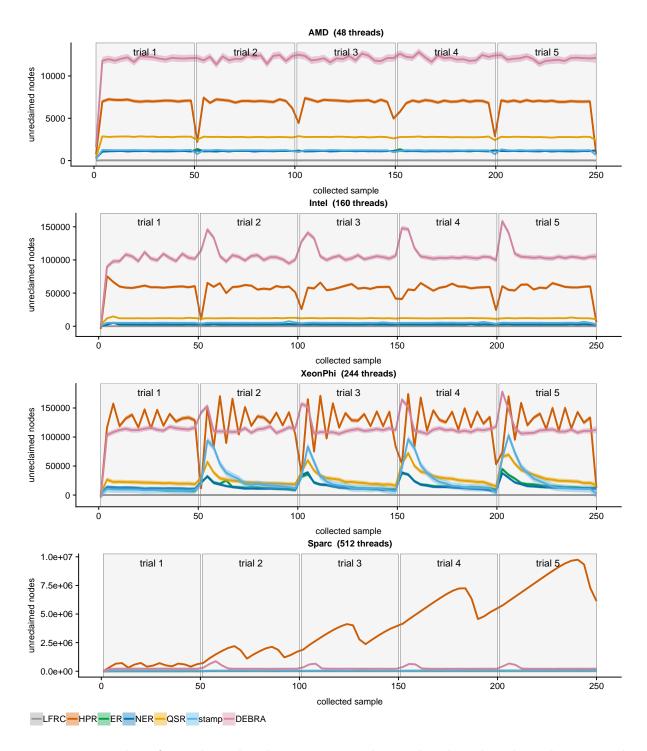


Figure 10: Number of unreclaimed nodes over time in the List benchmark with 10 elements and a workload of 80%.

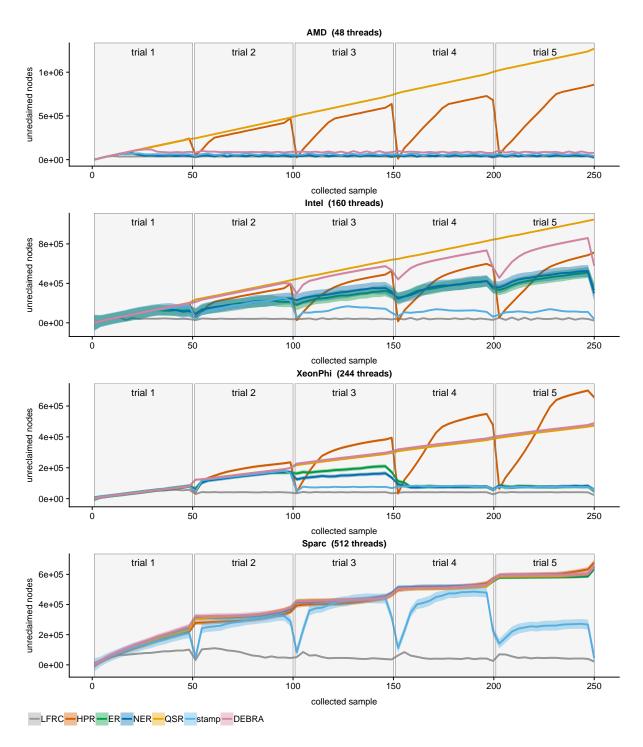


Figure 11: Number of unreclaimed nodes over time in the HashMap benchmark.

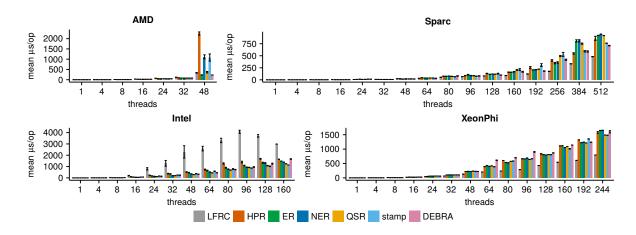


Figure 12: Queue benchmark with varying number of threads.

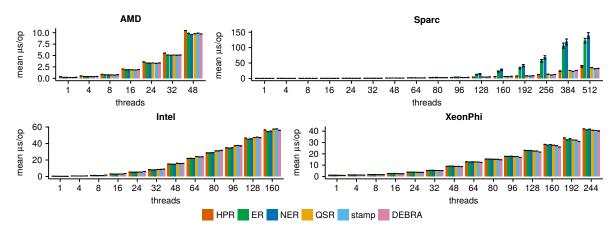


Figure 13: List benchmark with 10 elements, a workload of 20% and varying number of threads (without LFRC).

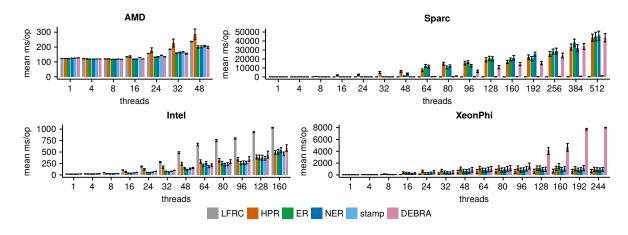


Figure 14: HashMap benchmark with varying number of threads.

runtime/operation is very similar for all schemes, in all experiments and on all machines, i.e., the impact of the memory manager is equally big/small for all schemes.

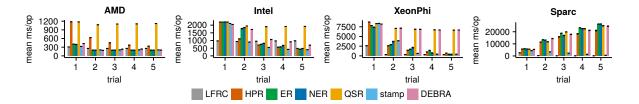


Figure 15: Development of runtime over time in the HashMap benchmark.

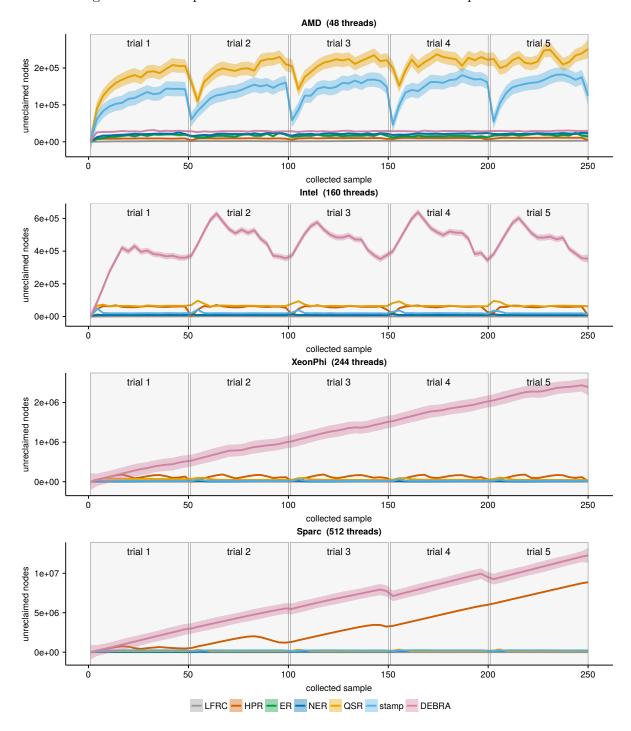


Figure 16: Number of unreclaimed of nodes over time in the Queue benchmark.

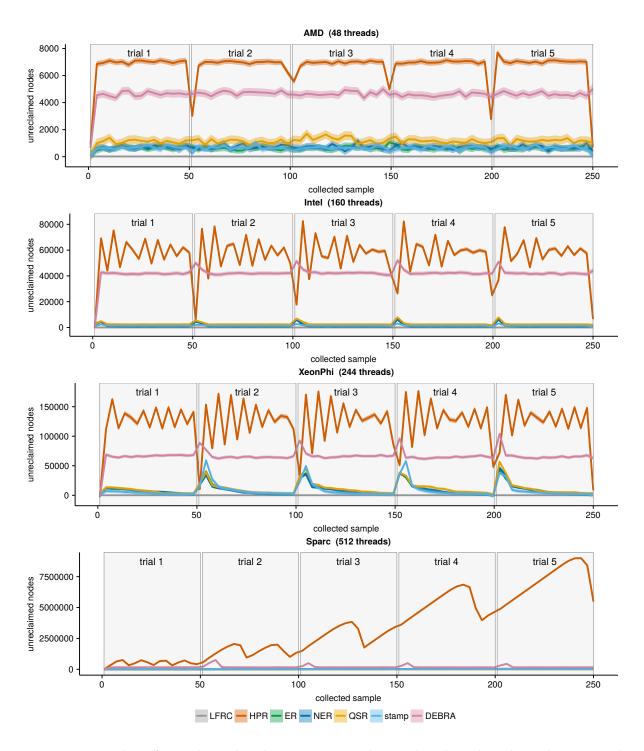


Figure 17: Number of unreclaimed nodes over time in the List benchmark with 10 elements and a workload of 20%.

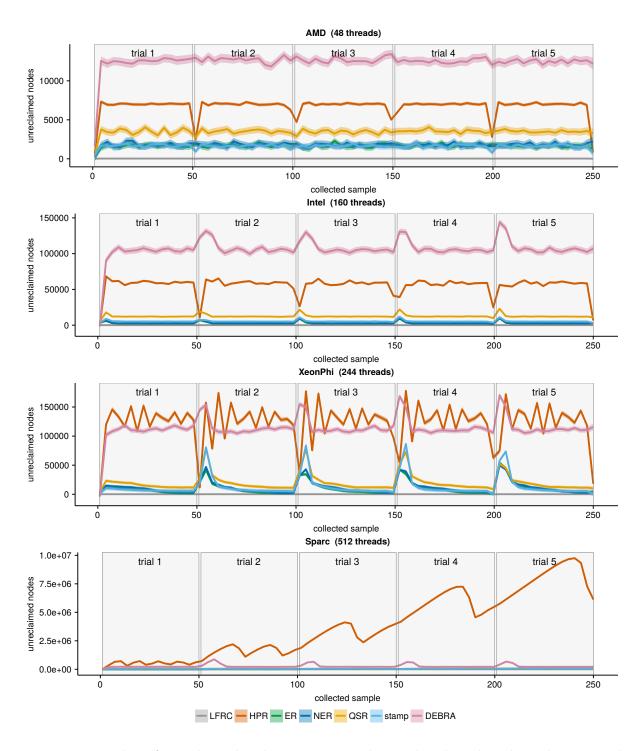


Figure 18: Number of unreclaimed nodes over time in the List benchmark with 10 elements and a workload of 80%.

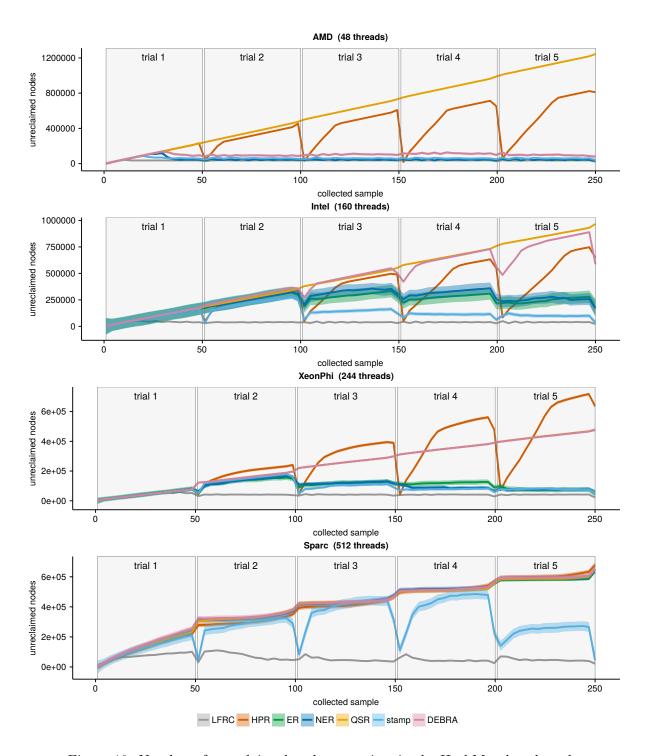


Figure 19: Number of unreclaimed nodes over time in the HashMap benchmark.